



INTRODUCCIÓN

Consultas con HIVE

- Trabajaremos con los dos archivos que vienen con la tarea (**movies.dat** -*películas con sus géneros*- y **ratings.dat** -*con un millón de votaciones*-) con la siguiente estructura:

movies.dat - id, película y sus géneros

```
SC13 > ANTONIO > DESKTOP > movies.dat
1*Toy Story (1995)*Animation|Children's|Comedy
2*Jumanji (1995)*Adventure|Children's|Fantasy
3*Grumpier Old Men (1995)*Comedy|Romance
4*Waiting to Exhale (1995)*Comedy|Drama
5*Father of the Bride Part II (1995)*Comedy
6*Heat (1995)*Action|Crime|Thriller
7*Sabrina (1995)*Comedy|Romance
```

rating.dat - id,usuario,voto,timestamp

```
1|1193|5|978300760
1|661|3|978302109
1|914|3|978301968
1|3408|4|978300275
1|2355|5|978824291
1|1197|3|978302268
```

- Para cada una de las consultas, mostrar su código y una captura con la salida de las mismas.

TABLES 5		🔍	➕	TABLE > USUARIOS		
		Search	🔍	COLUMNS	DDL	STORAGE INFORMATION
				AUTHORIZATION	DET	
	peliculas					
	votos			COLUMN NAME	COLUMN TYPE	CO
	usuarios			user_id	int	
	ml_items_managed			age	int	
	ml_user_info			gender	string	
				occupation	string	

```
[maria_dev@sandbox-hdp movielens]$ ls
movies.dat  u.data  u.data2.txt  u.item2          u.item2.txt  u.user2
ratings.dat u.data2 u.item      u.item2 con comas  u.user      u.user2.txt
[maria_dev@sandbox-hdp movielens]$
```

CONTENIDO

APARTADO A

1. Crea las tablas HIVE adecuadas para cargar los datos de cada uno de los archivos. Para mayor eficiencia, crearemos ambas tablas con 5 *buckets* sobre el **id** de la película. Tener especial cuidado en el campo que almacene el género de las películas.

Vamos a crear sobre la base de datos movielens las Tablas MOVIES Y RATINGS

```
hive> CREATE EXTERNAL TABLE IF NOT EXISTS MOVIES (
    >     movieId INT,
    >     title STRING,
    >     genres STRING
    > )
    > ROW FORMAT DELIMITED
    > FIELDS TERMINATED BY '|'
    > STORED AS TEXTFILE;
hive> LOAD DATA INPATH '/user/maria_dev/movielens/movies.dat'
    >         OVERWRITE INTO TABLE MOVIES;
```

Buket para movies

```
hive> CREATE TABLE IF NOT EXISTS MOVIES_bucketed (
    >     movieId INT,
    >     title STRING,
    >     genres STRING
    > )
    > CLUSTERED BY (movieId) INTO 5 BUCKETS
    > STORED AS ORC
    > TBLPROPERTIES ("orc.compress"="SNAPPY");
```

```
hive>
    > INSERT INTO TABLE MOVIES_bucketed
    > SELECT movieId, title, genres
    > FROM MOVIES;
```

VERTICES	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	SUCCEEDED	1	1	0	0	0	0
Reducer 2	SUCCEEDED	5	5	0	0	0	0

VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 36.42 s

Loading data to table movielens.movies_bucketed
Table movielens.movies_bucketed stats: [numFiles=5, numRows=3883, totalSize=84271, rawDataSize=783589]
OK

Creacion de tabla ratings

```
hive> CREATE EXTERNAL TABLE IF NOT EXISTS RATINGS (
    >     userId INT,
    >     movieId INT,
    >     rating DOUBLE,
    >     times_tamp BIGINT
    > )
    > ROW FORMAT DELIMITED
    > FIELDS TERMINATED BY '|'
    > STORED AS TEXTFILE;
hive> LOAD DATA INPATH '/user/maria_dev/movielens/ratings.dat'
    >         OVERWRITE INTO TABLE RATINGS;
```

Buket para ratings

```
hive> CREATE TABLE IF NOT EXISTS RATINGS_bucketed (
    >     userId INT,
    >     movieId INT,
    >     rating DOUBLE,
    >     times_tamp BIGINT
    > )
    > CLUSTERED BY (movieId) INTO 5 BUCKETS
    > STORED AS ORC
    > TBLPROPERTIES ("orc.compress"="SNAPPY");
hive> SET hive.enforce.bucketing = true;
hive>
    > INSERT INTO TABLE RATINGS_bucketed
    > SELECT userId, movieId, rating, times_tamp
    > FROM RATINGS;
Query ID = maria_dev_20251204094341_8a5e97e2-995c-48b3-8c8e-3420488f2e3c
Total jobs = 1
Launching Job 1 out of 1
Status: Running (Executing on YARN cluster with App id application_1764836683577_0004)
```

VERTICES	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	SUCCEEDED	2	2	0	0	0	0
Reducer 2	SUCCEEDED	5	5	0	0	0	0
VERTICES: 02/02 [=====>>] 100% ELAPSED TIME: 12.32 s							
Loading data to table movielens.ratings_bucketed							
Table movielens.ratings_bucketed stats: [numFiles=5, numRows=1000209, totalSize=8497198, rawDataSize=24005016]							
OK							
Time taken: 14.838 seconds							

2. Mostrar las consultas de creación y carga de las tablas con una captura que muestre los datos cargados.

```
hive> SELECT * FROM MOVIES LIMIT 10;
OK
1      Toy Story (1995)      Animation
2      Jumanji (1995)       Adventure
3      Grumpier Old Men (1995) Comedy
4      Waiting to Exhale (1995)    Comedy
5      Father of the Bride Part II (1995)   Comedy
6      Heat (1995)          Action
7      Sabrina (1995)        Comedy
8      Tom and Huck (1995)     Adventure
9      Sudden Death (1995)     Action
10     GoldenEye (1995)       Action
Time taken: 0.153 seconds, Fetched: 10 row(s)
hive> SELECT * FROM RATINGS LIMIT 10;
OK
1      1193      5.0      978300760
1      661       3.0      978302109
1      914       3.0      978301968
1      3408      4.0      978300275
1      2355      5.0      978824291
1      1197      3.0      978302268
1      1287      5.0      978302039
1      2804      5.0      978300719
1      594       4.0      978302268
1      919       4.0      978301368
Time taken: 0.147 seconds, Fetched: 10 row(s)
hive>
```

3. Mostrar una captura de las tablas en el *warehouse* de HIVE donde se vean los *buckets*.

```

hive> DESCRIBE FORMATTED MOVIES;
OK
# col_name          data_type          comment
movieid            int
title              string
genres             string

# Detailed Table Information
Database:          movielens
Owner:              maria_dev
CreateTime:         Fri Dec 05 20:03:28 UTC 2025
LastAccessTime:    UNKNOWN
Protect Mode:      None
Retention:          0
Location:          hdfs://sandbox-hdp.hortonworks.com:8020/apps/hive/warehouse/movielens.db/movies
Table Type:        EXTERNAL_TABLE
Table Parameters:
  EXTERNAL           TRUE
  numFiles           1
  numRows            0
  rawDataSize        0
  totalSize          163542
  transient_lastDdlTime 1764965051

# Storage Information
SerDe Library:     org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe
InputFormat:        org.apache.hadoop.mapred.TextInputFormat
OutputFormat:       org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat
Compressed:         No
Num Buckets:       -1
Bucket Columns:    []
Sort Columns:       []
Storage Desc Params:
  field.delim      |
  serialization.format  |
Time taken: 0.512 seconds, Fetched: 34 row(s)

```

```

hive> DESCRIBE FORMATTED MOVIES_bucket;
FAILED: SemanticException [Error 10001]: Table not found MOVIES_bucket
hive> DESCRIBE FORMATTED MOVIES_bucketed;
OK
# col_name          data_type          comment
movieid            int
title              string
genres             string

# Detailed Table Information
Database:          movielens
Owner:              maria_dev
CreateTime:         Fri Dec 05 20:07:23 UTC 2025
LastAccessTime:    UNKNOWN
Protect Mode:      None
Retention:          0
Location:          hdfs://sandbox-hdp.hortonworks.com:8020/apps/hive/warehouse/movielens.db/movies_bucketed
Table Type:        MANAGED_TABLE
Table Parameters:
  COLUMN_STATS_ACCURATE  {"BASIC_STATS": "true"}
  numFiles               5
  numRows                3883
  orc.compress           SNAPPY
  rawDataSize            783589
  totalSize               84271
  transient_lastDdlTime  1764965399

# Storage Information
SerDe Library:     org.apache.hadoop.hive.ql.io.orc.OrcSerde
InputFormat:        org.apache.hadoop.hive.ql.io.orc.OrcInputFormat
OutputFormat:       org.apache.hadoop.hive.ql.io.orc.OrcOutputFormat
Compressed:         No
Num Buckets:       5
Bucket Columns:    [movieid]
Sort Columns:       []
Storage Desc Params:
  serialization.format  1
Time taken: 0.523 seconds, Fetched: 34 row(s)

```

```
hive> DESCRIBE FORMATTED RATINGS;
OK
# col_name          data_type         comment
userid              int
movieid             int
rating              double
times_tamp          bigint

# Detailed Table Information
Database:          movielens
Owner:              maria_dev
CreateTime:         Thu Dec 04 09:30:48 UTC 2025
LastAccessTime:    UNKNOWN
Protect Mode:      None
Retention:          0
Location:          hdfs://sandbox-hdp.hortonworks.com:8020/apps/hive/warehouse/movielens.db/ratings
Table Type:        EXTERNAL_TABLE
Table Parameters:
  EXTERNAL           TRUE
  numFiles           1
  numRows            0
  rawDataSize       0
  totalSize          21593504
  transient_lastDdlTime 1764840701

# Storage Information
SerDe Library:     org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe
InputFormat:        org.apache.hadoop.mapred.TextInputFormat
OutputFormat:       org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat
Compressed:         No
Num Buckets:       -1
Bucket Columns:    []
Sort Columns:       []
Storage Desc Params:
  field.delim          |
  serialization.format  |

Time taken: 0.54 seconds, Fetched: 35 row(s)
hive>
```

```

hive> DESCRIBE FORMATTED RATINGS_bucketed;
OK
# col_name          data_type        comment
userid              int
movieid             int
rating              double
times_tamp          bigint

# Detailed Table Information
Database:          movielens
Owner:              maria_dev
CreateTime:         Thu Dec 04 09:43:17 UTC 2025
LastAccessTime:    UNKNOWN
Protect Mode:      None
Retention:         0
Location:          hdfs://sandbox-hdp.hortonworks.com:8020/apps/hive/warehouse/movielens.d
Table Type:        MANAGED_TABLE
Table Parameters:
  COLUMN_STATS_ACCURATE  {"BASIC_STATS": "true"}
  numFiles                5
  numRows                 1000209
  orc.compress            SNAPPY
  rawDataSize             24005016
  totalSize                8497198
  transient_lastDdlTime   1764841435

# Storage Information
SerDe Library:     org.apache.hadoop.hive.ql.io.orc.OrcSerde
InputFormat:        org.apache.hadoop.hive.ql.io.orc.OrcInputFormat
OutputFormat:       org.apache.hadoop.hive.ql.io.orc.OrcOutputFormat
Compressed:         No
Num Buckets:       5
Bucket Columns:    [movieid]
Sort Columns:       []
Storage Desc Params:
  serialization.format    1
Time taken: 0.554 seconds, Fetched: 35 row(s)

```

CONTENIDO

APARTADO B

1. Mediante una consulta en HIVE, encontrar las cinco películas (código, título y media de votos) mejor valoradas, que hayan sido votadas al menos por 10 usuarios.

```

hive> SELECT
    >     m.movieId,
    >     m.title,
    >     ROUND(avg_r.avg_rating, 3) AS avg_rating,
    >     avg_r.cnt AS votes
    > FROM
    >     (SELECT movieId, AVG(rating) AS avg_rating, COUNT(*) AS cnt
    >      FROM movielens.ratings_bucketed
    >      GROUP BY movieId
    >      HAVING cnt >= 10
    > ) avg_r
    > JOIN movielens.movies_bucketed m
    >   ON m.movieId = avg_r.movieId
    > ORDER BY avg_rating DESC, votes DESC
    > LIMIT 5;

```

VERTICES	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	SUCCEEDED	1	1	0	0	0	0
Map 4	SUCCEEDED	1	1	0	0	0	0
Reducer 2	SUCCEEDED	1	1	0	0	0	0
Reducer 3	SUCCEEDED	1	1	0	0	0	0
VERTICES: 04/04 [=====>>] 100% ELAPSED TIME: 12.76 s							
OK							
2905	Sanjuro (1962)	4.609	69				
2019	Seven Samurai (The Magnificent Seven) (Shichinin no samurai) (1954)			4.561	628		
318	Shawshank Redemption, The (1994)	4.555	2227				
858	Godfather, The (1972)	4.525	2223				
745	Close Shave, A (1995)	4.521	657				
Time taken: 14.859 seconds, Fetched: 5 row(s)							

CONTENIDO

APARTADO C

1. Encontrar las cinco películas más antiguas con una valoración media por encima de 4 puntos.

```

hive> SELECT
    >     t.movieId,
    >     t.title,
    >     t.year,
    >     ROUND(t.avg_rating, 3) AS avg_rating
    > FROM (
    >     SELECT
    >         m.movieId,
    >         m.title,
    >         CAST(regexp_extract(m.title, '\\\\((\\\\d{4})\\\\)$', 1) AS INT) AS year,
    >         AVG(r.rating) AS avg_rating
    >     FROM movielens.movies_bucketed m
    >     JOIN movielens.ratings_bucketed r ON m.movieId = r.movieId
    >     GROUP BY m.movieId, m.title
    > ) t
    > WHERE t.year IS NOT NULL
    >     AND t.avg_rating > 4
    > ORDER BY t.year ASC
    > LIMIT 5.

```

	VERTICES	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	SUCCEEDED	1	1	0	0	0	0	0
Map 2	SUCCEEDED	1	1	0	0	0	0	0
Reducer 3	SUCCEEDED	1	1	0	0	0	0	0
Reducer 4	SUCCEEDED	1	1	0	0	0	0	0
<hr/>								
VERTICES: 04/04 [=====>>] 100% ELAPSED TIME: 12.08 s								
<hr/>								
OK								
3629	Gold Rush, The (1925)	1925	4.189					
2010	Metropolis (1926)	1926	4.082					
3517	Bells, The (1926)	1926	4.5					
3022	General, The (1927)	1927	4.369					
1927	All Quiet on the Western Front (1930)	1930	4.194					
Time taken: 13.362 seconds, Fetched: 5 row(s)								

CONTENIDO

APARTADO D

Investigando las funciones que nos ofrece HIVE para el manejo de cadenas....

1. Muestra los cinco años en que se editaron más películas indicando el número de ellas para cada año.



Big Data

```
hive> SELECT
    >     year,
    >     COUNT(*) AS movies_count
    > FROM (
    >     SELECT
    >         movieId,
    >         CAST(regexp_extract(title, '\\((\\d{4})\\)\\$', 1) AS INT) AS year
    >     FROM movielens.movies_bucketed
    > ) x
    > WHERE year IS NOT NULL
    > GROUP BY year
    > ORDER BY movies_count DESC
    > LIMIT 5;
```

VERTICES	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	SUCCEEDED	1	1	0	0	0	0
Reducer 2	SUCCEEDED	1	1	0	0	0	0
Reducer 3	SUCCEEDED	1	1	0	0	0	0

VERTICES: 03/03 [=====>>] 100% ELAPSED TIME: 7.80 s

OK
1996 345
1995 342
1998 337
1997 314
1999 283

CONTENIDO

APARTADO E

Consultando el enlace de abajo:

<https://www.bigdatainrealworld.com/what-is-the-difference-between-explode-and-lateralview-explode-in-hive/>

1. Muestra los diez géneros de películas más frecuentes.

```
hive> SELECT
    >     genre,
    >     COUNT(*) AS freq
    > FROM movielens.movies_bucketed
    > LATERAL VIEW explode(split(genres, '|')) g AS genre
    > GROUP BY genre
    > ORDER BY freq DESC
    > LIMIT 10;
```

VERTICES	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 1	SUCCEEDED	1	1	0	0	0	0
Reducer 2	SUCCEEDED	1	1	0	0	0	0
Reducer 3	SUCCEEDED	1	1	0	0	0	0
VERTICES: 03/03 [=====>>] 100% ELAPSED TIME: 8.34 s							
OK							
Drama	1176						
Comedy	1022						
Action	503						
Horror	262						
Adventure	155						
Crime	131						
Documentary	123						
Thriller	101						
Animation	90						
Children's	89						

2. Partiendo de la consulta del Apartado B de arriba. ¿Qué géneros son los más frecuentes en dichas películas?

```
hive> WITH top5 AS (
    >   SELECT
    >     m.movieId
    >   FROM (
    >     SELECT movieId, AVG(rating) AS avg_rating, COUNT(*) AS cnt
    >       FROM movielens.ratings_bucketed
    >     GROUP BY movieId
    >     HAVING cnt >= 10
    >     ORDER BY avg_rating DESC, cnt DESC
    >     LIMIT 5
    >   ) t
    >   JOIN movielens.movies_bucketed m ON t.movieId = m.movieId
    > )
    >
    > SELECT
    >   genre,
    >   COUNT(*) AS freq_in_top5
    > FROM top5 t
    > JOIN movielens.movies_bucketed m ON t.movieId = m.movieId
    > LATERAL VIEW explode(split(m.genres, '\\|')) g AS genre
    > GROUP BY genre
    > ORDER BY freq_in_top5 DESC;
```

